

stylegan

April 1, 2024

0.1 Setup

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[ ]: !pip install tensorflow_addons
```

Collecting tensorflow_addons

Downloading tensorflow_addons-0.23.0-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (611 kB)

611.8/611.8

kB 3.5 MB/s eta 0:00:00

Requirement already satisfied: packaging in

/usr/local/lib/python3.10/dist-packages (from tensorflow_addons) (24.0)

Collecting typeguard<3.0.0,>=2.7 (from tensorflow_addons)

Downloading typeguard-2.13.3-py3-none-any.whl (17 kB)

Installing collected packages: typeguard, tensorflow_addons

Successfully installed tensorflow_addons-0.23.0 typeguard-2.13.3

```
[ ]: import os
import numpy as np
import matplotlib.pyplot as plt

from functools import partial

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential
from tensorflow_addons.layers import InstanceNormalization

import gdown
from zipfile import ZipFile
```

/usr/local/lib/python3.10/dist-packages/tensorflow_addons/utils/tfa_eol_msg.py:23: UserWarning:

TensorFlow Addons (TFA) has ended development and introduction of new features. TFA has entered a minimal maintenance and release mode until a planned end of life in May 2024.

Please modify downstream libraries to take dependencies from other repositories in our TensorFlow community (e.g. Keras, Keras-CV, and Keras-NLP).

For more information see: <https://github.com/tensorflow/addons/issues/2807>

```
warnings.warn(
```

```
[ ]: def log2(x):
    return int(np.log2(x))

# we use different batch size for different resolution, so larger image size
# could fit into GPU memory. The keys is image resolution in log2
batch_sizes = {2: 16, 3: 16, 4: 16, 5: 16, 6: 16, 7: 8, 8: 4, 9: 2, 10: 1}
# We adjust the train step accordingly
train_step_ratio = {k: batch_sizes[2] / v for k, v in batch_sizes.items()}

os.makedirs("celeba_gan")

url = "https://drive.google.com/uc?id=107m1010EJjLE5QxLZiM9Fpjs70j6e684"
output = "celeba_gan/data.zip"
gdown.download(url, output, quiet=True)

with ZipFile("celeba_gan/data.zip", "r") as zipobj:
    zipobj.extractall("celeba_gan")

# Create a dataset from our folder, and rescale the images to the [0-1] range:

ds_train = keras.utils.image_dataset_from_directory(
    "celeba_gan", label_mode=None, image_size=(64, 64), batch_size=32
)

def resize_image(res, image):
    # only downsampling, so use nearest neighbor that is faster to run
    image = tf.image.resize(
        image, (res, res), method=tf.image.ResizeMethod.NEAREST_NEIGHBOR
    )
    image = tf.cast(image, tf.float32) / 127.5 - 1.0
    return image
```

```
def create_dataloader(res):
    batch_size = batch_sizes[log2(res)]
    dl = ds_train.map(partial(resize_image, res), num_parallel_calls=tf.data.
    ↪AUTOTUNE).unbatch()
    dl = dl.shuffle(200).batch(batch_size, drop_remainder=True).prefetch(1).
    ↪repeat()
    return dl
```

Found 202599 files belonging to 1 classes.

```
[ ]: def plot_images(images, log2_res, fname=""):
    scales = {2: 0.5, 3: 1, 4: 2, 5: 3, 6: 4, 7: 5, 8: 6, 9: 7, 10: 8}
    scale = scales[log2_res]

    grid_col = min(images.shape[0], int(32 // scale))
    grid_row = 1

    f, axarr = plt.subplots(
        grid_row, grid_col, figsize=(grid_col * scale, grid_row * scale)
    )

    for row in range(grid_row):
        ax = axarr if grid_row == 1 else axarr[row]
        for col in range(grid_col):
            ax[col].imshow(images[row * grid_col + col])
            ax[col].axis("off")
    plt.show()
    if fname:
        f.savefig(fname)
```

```
[ ]: def fade_in(alpha, a, b):
    return alpha * a + (1.0 - alpha) * b

def wasserstein_loss(y_true, y_pred):
    return -tf.reduce_mean(y_true * y_pred)

def pixel_norm(x, epsilon=1e-8):
    return x / tf.math.sqrt(tf.reduce_mean(x ** 2, axis=-1, keepdims=True) +
    ↪epsilon)

def minibatch_std(input_tensor, epsilon=1e-8):
    n, h, w, c = tf.shape(input_tensor)
    group_size = tf.minimum(4, n)
```

```

x = tf.reshape(input_tensor, [group_size, -1, h, w, c])
group_mean, group_var = tf.nn.moments(x, axes=(0), keepdims=False)
group_std = tf.sqrt(group_var + epsilon)
avg_std = tf.reduce_mean(group_std, axis=[1, 2, 3], keepdims=True)
x = tf.tile(avg_std, [group_size, h, w, 1])
return tf.concat([input_tensor, x], axis=-1)

class EqualizedConv(layers.Layer):
    def __init__(self, out_channels, kernel=3, gain=2, **kwargs):
        super().__init__(**kwargs)
        self.kernel = kernel
        self.out_channels = out_channels
        self.gain = gain
        self.pad = kernel != 1

    def build(self, input_shape):
        self.in_channels = input_shape[-1]
        initializer = keras.initializers.RandomNormal(mean=0.0, stddev=1.0)
        self.w = self.add_weight(
            shape=[self.kernel, self.kernel, self.in_channels, self.
↪out_channels],
            initializer=initializer,
            trainable=True,
            name="kernel",
        )
        self.b = self.add_weight(
            shape=(self.out_channels,), initializer="zeros", trainable=True,
↪name="bias"
        )
        fan_in = self.kernel * self.kernel * self.in_channels
        self.scale = tf.sqrt(self.gain / fan_in)

    def call(self, inputs):
        if self.pad:
            x = tf.pad(inputs, [[0, 0], [1, 1], [1, 1], [0, 0]], mode="REFLECT")
        else:
            x = inputs
        output = (
            tf.nn.conv2d(x, self.scale * self.w, strides=1, padding="VALID") +
↪self.b
        )
        return output

class EqualizedDense(layers.Layer):
    def __init__(self, units, gain=2, learning_rate_multiplier=1, **kwargs):

```

```

    super().__init__(**kwargs)
    self.units = units
    self.gain = gain
    self.learning_rate_multiplier = learning_rate_multiplier

    def build(self, input_shape):
        self.in_channels = input_shape[-1]
        initializer = keras.initializers.RandomNormal(
            mean=0.0, stddev=1.0 / self.learning_rate_multiplier
        )
        self.w = self.add_weight(
            shape=[self.in_channels, self.units],
            initializer=initializer,
            trainable=True,
            name="kernel",
        )
        self.b = self.add_weight(
            shape=(self.units,), initializer="zeros", trainable=True,
↪name="bias"
        )
        fan_in = self.in_channels
        self.scale = tf.sqrt(self.gain / fan_in)

    def call(self, inputs):
        output = tf.add(tf.matmul(inputs, self.scale * self.w), self.b)
        return output * self.learning_rate_multiplier

class AddNoise(layers.Layer):
    def build(self, input_shape):
        n, h, w, c = input_shape[0]
        initializer = keras.initializers.RandomNormal(mean=0.0, stddev=1.0)
        self.b = self.add_weight(
            shape=[1, 1, 1, c], initializer=initializer, trainable=True,
↪name="kernel"
        )

    def call(self, inputs):
        x, noise = inputs
        output = x + self.b * noise
        return output

class AdaIN(layers.Layer):
    def __init__(self, gain=1, **kwargs):
        super().__init__(**kwargs)
        self.gain = gain

```

```

def build(self, input_shapes):
    x_shape = input_shapes[0]
    w_shape = input_shapes[1]

    self.w_channels = w_shape[-1]
    self.x_channels = x_shape[-1]

    self.dense_1 = EqualizedDense(self.x_channels, gain=1)
    self.dense_2 = EqualizedDense(self.x_channels, gain=1)

def call(self, inputs):
    x, w = inputs
    ys = tf.reshape(self.dense_1(w), (-1, 1, 1, self.x_channels))
    yb = tf.reshape(self.dense_2(w), (-1, 1, 1, self.x_channels))
    return ys * x + yb

```

```

[ ]: def Mapping(num_stages, input_shape=512):
    z = layers.Input(shape=(input_shape))
    w = pixel_norm(z)
    for i in range(8):
        w = EqualizedDense(512, learning_rate_multiplier=0.01)(w)
        w = layers.LeakyReLU(0.2)(w)
    w = tf.tile(tf.expand_dims(w, 1), (1, num_stages, 1))
    return keras.Model(z, w, name="mapping")

class Generator:
    def __init__(self, start_res_log2, target_res_log2):
        self.start_res_log2 = start_res_log2
        self.target_res_log2 = target_res_log2
        self.num_stages = target_res_log2 - start_res_log2 + 1
        # list of generator blocks at increasing resolution
        self.g_blocks = []
        # list of layers to convert g_block activation to RGB
        self.to_rgb = []
        # list of noise input of different resolutions into g_blocks
        self.noise_inputs = []
        # filter size to use at each stage, keys are log2(resolution)
        self.filter_nums = {
            0: 512,
            1: 512,
            2: 512, # 4x4
            3: 512, # 8x8
            4: 512, # 16x16
            5: 512, # 32x32
            6: 256, # 64x64

```

```

        7: 128, # 128x128
        8: 64, # 256x256
        9: 32, # 512x512
        10: 16,
    } # 1024x1024

    start_res = 2 ** start_res_log2
    self.input_shape = (start_res, start_res, self.
↪filter_nums[start_res_log2])
    self.g_input = layers.Input(self.input_shape, name="generator_input")

    for i in range(start_res_log2, target_res_log2 + 1):
        filter_num = self.filter_nums[i]
        res = 2 ** i
        self.noise_inputs.append(
            layers.Input(shape=(res, res, 1), name=f"noise_{res}x{res}")
        )
        to_rgb = Sequential(
            [
                layers.InputLayer(input_shape=(res, res, filter_num)),
                EqualizedConv(3, 1, gain=1),
            ],
            name=f"to_rgb_{res}x{res}",
        )
        self.to_rgb.append(to_rgb)
        is_base = i == self.start_res_log2
        if is_base:
            input_shape = (res, res, self.filter_nums[i - 1])
        else:
            input_shape = (2 ** (i - 1), 2 ** (i - 1), self.filter_nums[i - 1])
↪1])

        g_block = self.build_block(
            filter_num, res=res, input_shape=input_shape, is_base=is_base
        )
        self.g_blocks.append(g_block)

    def build_block(self, filter_num, res, input_shape, is_base):
        input_tensor = layers.Input(shape=input_shape, name=f"g_{res}")
        noise = layers.Input(shape=(res, res, 1), name=f"noise_{res}")
        w = layers.Input(shape=512)
        x = input_tensor

        if not is_base:
            x = layers.UpSampling2D((2, 2))(x)
            x = EqualizedConv(filter_num, 3)(x)

        x = AddNoise()([x, noise])

```

```

x = layers.LeakyReLU(0.2)(x)
x = InstanceNormalization()(x)
x = AdaIN()([x, w])

x = EqualizedConv(filter_num, 3)(x)
x = AddNoise()([x, noise])
x = layers.LeakyReLU(0.2)(x)
x = InstanceNormalization()(x)
x = AdaIN()([x, w])
return keras.Model([input_tensor, w, noise], x,
name=f"genblock_{res}x{res}")

def grow(self, res_log2):
    res = 2 ** res_log2

    num_stages = res_log2 - self.start_res_log2 + 1
    w = layers.Input(shape=(self.num_stages, 512), name="w")

    alpha = layers.Input(shape=(1), name="g_alpha")
    x = self.g_blocks[0]([self.g_input, w[:, 0], self.noise_inputs[0]])

    if num_stages == 1:
        rgb = self.to_rgb[0](x)
    else:
        for i in range(1, num_stages - 1):

            x = self.g_blocks[i]([x, w[:, i], self.noise_inputs[i]])

        old_rgb = self.to_rgb[num_stages - 2](x)
        old_rgb = layers.UpSampling2D((2, 2))(old_rgb)

        i = num_stages - 1
        x = self.g_blocks[i]([x, w[:, i], self.noise_inputs[i]])

        new_rgb = self.to_rgb[i](x)

        rgb = fade_in(alpha[0], new_rgb, old_rgb)

    return keras.Model(
        [self.g_input, w, self.noise_inputs, alpha],
        rgb,
        name=f"generator_{res}_x_{res}",
    )

class Discriminator:
    def __init__(self, start_res_log2, target_res_log2):

```



```

self.start_res_log2 = start_res_log2
self.target_res_log2 = target_res_log2
self.num_stages = target_res_log2 - start_res_log2 + 1
# filter size to use at each stage, keys are log2(resolution)
self.filter_nums = {
    0: 512,
    1: 512,
    2: 512, # 4x4
    3: 512, # 8x8
    4: 512, # 16x16
    5: 512, # 32x32
    6: 256, # 64x64
    7: 128, # 128x128
    8: 64, # 256x256
    9: 32, # 512x512
    10: 16,
} # 1024x1024
# list of discriminator blocks at increasing resolution
self.d_blocks = []
# list of layers to convert RGB into activation for d_blocks inputs
self.from_rgb = []

for res_log2 in range(self.start_res_log2, self.target_res_log2 + 1):
    res = 2 ** res_log2
    filter_num = self.filter_nums[res_log2]
    from_rgb = Sequential(
        [
            layers.InputLayer(
                input_shape=(res, res, 3), name=f"from_rgb_input_{res}"
            ),
            EqualizedConv(filter_num, 1),
            layers.LeakyReLU(0.2),
        ],
        name=f"from_rgb_{res}",
    )

    self.from_rgb.append(from_rgb)

    input_shape = (res, res, filter_num)
    if len(self.d_blocks) == 0:
        d_block = self.build_base(filter_num, res)
    else:
        d_block = self.build_block(
            filter_num, self.filter_nums[res_log2 - 1], res
        )

    self.d_blocks.append(d_block)

```

```

def build_base(self, filter_num, res):
    input_tensor = layers.Input(shape=(res, res, filter_num),
    ↪name=f"d_{res}")
    x = minibatch_std(input_tensor)
    x = EqualizedConv(filter_num, 3)(x)
    x = layers.LeakyReLU(0.2)(x)
    x = layers.Flatten()(x)
    x = EqualizedDense(filter_num)(x)
    x = layers.LeakyReLU(0.2)(x)
    x = EqualizedDense(1)(x)
    return keras.Model(input_tensor, x, name=f"d_{res}")

def build_block(self, filter_num_1, filter_num_2, res):
    input_tensor = layers.Input(shape=(res, res, filter_num_1),
    ↪name=f"d_{res}")
    x = EqualizedConv(filter_num_1, 3)(input_tensor)
    x = layers.LeakyReLU(0.2)(x)
    x = EqualizedConv(filter_num_2)(x)
    x = layers.LeakyReLU(0.2)(x)
    x = layers.AveragePooling2D((2, 2))(x)
    return keras.Model(input_tensor, x, name=f"d_{res}")

def grow(self, res_log2):
    res = 2 ** res_log2
    idx = res_log2 - self.start_res_log2
    alpha = layers.Input(shape=(1), name="d_alpha")
    input_image = layers.Input(shape=(res, res, 3), name="input_image")
    x = self.from_rgb[idx](input_image)
    x = self.d_blocks[idx](x)
    if idx > 0:
        idx -= 1
        downsized_image = layers.AveragePooling2D((2, 2))(input_image)
        y = self.from_rgb[idx](downsized_image)
        x = fade_in(alpha[0], x, y)

        for i in range(idx, -1, -1):
            x = self.d_blocks[i](x)
    return keras.Model([input_image, alpha], x,
    ↪name=f"discriminator_{res}_x_{res}")

```

```

[ ]: class StyleGAN(tf.keras.Model):
    def __init__(self, z_dim=512, target_res=64, start_res=4):
        super().__init__()
        self.z_dim = z_dim

        self.target_res_log2 = log2(target_res)

```

```

self.start_res_log2 = log2(start_res)
self.current_res_log2 = self.target_res_log2
self.num_stages = self.target_res_log2 - self.start_res_log2 + 1

self.alpha = tf.Variable(1.0, dtype=tf.float32, trainable=False,
↪name="alpha")

self.mapping = Mapping(num_stages=self.num_stages)
self.d_builder = Discriminator(self.start_res_log2, self.
↪target_res_log2)
self.g_builder = Generator(self.start_res_log2, self.target_res_log2)
self.g_input_shape = self.g_builder.input_shape

self.phase = None
self.train_step_counter = tf.Variable(0, dtype=tf.int32,
↪trainable=False)

self.loss_weights = {"gradient_penalty": 10, "drift": 0.001}

def grow_model(self, res):
    tf.keras.backend.clear_session()
    res_log2 = log2(res)
    self.generator = self.g_builder.grow(res_log2)
    self.discriminator = self.d_builder.grow(res_log2)
    self.current_res_log2 = res_log2
    print(f"\nModel resolution:{res}x{res}")

def compile(
    self, steps_per_epoch, phase, res, d_optimizer, g_optimizer, *args,
↪**kwargs
):
    self.loss_weights = kwargs.pop("loss_weights", self.loss_weights)
    self.steps_per_epoch = steps_per_epoch
    if res != 2 ** self.current_res_log2:
        self.grow_model(res)
        self.d_optimizer = d_optimizer
        self.g_optimizer = g_optimizer

    self.train_step_counter.assign(0)
    self.phase = phase
    self.d_loss_metric = keras.metrics.Mean(name="d_loss")
    self.g_loss_metric = keras.metrics.Mean(name="g_loss")
    super().compile(*args, **kwargs)

@property
def metrics(self):
    return [self.d_loss_metric, self.g_loss_metric]

```

```

def generate_noise(self, batch_size):
    noise = [
        tf.random.normal((batch_size, 2 ** res, 2 ** res, 1))
        for res in range(self.start_res_log2, self.target_res_log2 + 1)
    ]
    return noise

def gradient_loss(self, grad):
    loss = tf.square(grad)
    loss = tf.reduce_sum(loss, axis=tf.range(1, tf.size(tf.shape(loss))))
    loss = tf.sqrt(loss)
    loss = tf.reduce_mean(tf.square(loss - 1))
    return loss

def train_step(self, real_images):

    self.train_step_counter.assign_add(1)

    if self.phase == "TRANSITION":
        self.alpha.assign(
            tf.cast(self.train_step_counter / self.steps_per_epoch, tf.
↪float32)
        )
    elif self.phase == "STABLE":
        self.alpha.assign(1.0)
    else:
        raise NotImplementedError

    alpha = tf.expand_dims(self.alpha, 0)
    batch_size = tf.shape(real_images)[0]
    real_labels = tf.ones(batch_size)
    fake_labels = -tf.ones(batch_size)

    z = tf.random.normal((batch_size, self.z_dim))
    const_input = tf.ones(tuple([batch_size] + list(self.g_input_shape)))
    noise = self.generate_noise(batch_size)

    # generator
    with tf.GradientTape() as g_tape:
        w = self.mapping(z)
        fake_images = self.generator([const_input, w, noise, alpha])
        pred_fake = self.discriminator([fake_images, alpha])
        g_loss = wasserstein_loss(real_labels, pred_fake)

        trainable_weights = (
            self.mapping.trainable_weights + self.generator.
↪trainable_weights

```

```

    )
    gradients = g_tape.gradient(g_loss, trainable_weights)
    self.g_optimizer.apply_gradients(zip(gradients, trainable_weights))

    # discriminator
    with tf.GradientTape() as gradient_tape, tf.GradientTape() as
↪total_tape:
        # forward pass
        pred_fake = self.discriminator([fake_images, alpha])
        pred_real = self.discriminator([real_images, alpha])

        epsilon = tf.random.uniform((batch_size, 1, 1, 1))
        interpolates = epsilon * real_images + (1 - epsilon) * fake_images
        gradient_tape.watch(interpolates)
        pred_fake_grad = self.discriminator([interpolates, alpha])

        # calculate losses
        loss_fake = wasserstein_loss(fake_labels, pred_fake)
        loss_real = wasserstein_loss(real_labels, pred_real)
        loss_fake_grad = wasserstein_loss(fake_labels, pred_fake_grad)

        # gradient penalty
        gradients_fake = gradient_tape.gradient(loss_fake_grad,
↪[interpolates])
        gradient_penalty = self.loss_weights[
            "gradient_penalty"
        ] * self.gradient_loss(gradients_fake)

        # drift loss
        all_pred = tf.concat([pred_fake, pred_real], axis=0)
        drift_loss = self.loss_weights["drift"] * tf.reduce_mean(all_pred,
↪** 2)

        d_loss = loss_fake + loss_real + gradient_penalty + drift_loss

        gradients = total_tape.gradient(
            d_loss, self.discriminator.trainable_weights
        )
        self.d_optimizer.apply_gradients(
            zip(gradients, self.discriminator.trainable_weights)
        )

    # Update metrics
    self.d_loss_metric.update_state(d_loss)
    self.g_loss_metric.update_state(g_loss)
    return {
        "d_loss": self.d_loss_metric.result(),

```

```

        "g_loss": self.g_loss_metric.result(),
    }

    def call(self, inputs: dict()):
        style_code = inputs.get("style_code", None)
        z = inputs.get("z", None)
        noise = inputs.get("noise", None)
        batch_size = inputs.get("batch_size", 1)
        alpha = inputs.get("alpha", 1.0)
        alpha = tf.expand_dims(alpha, 0)
        if style_code is None:
            if z is None:
                z = tf.random.normal((batch_size, self.z_dim))
                style_code = self.mapping(z)

        if noise is None:
            noise = self.generate_noise(batch_size)

        # self.alpha.assign(alpha)

        const_input = tf.ones(tuple([batch_size] + list(self.g_input_shape)))
        images = self.generator([const_input, style_code, noise, alpha])
        images = np.clip((images * 0.5 + 0.5) * 255, 0, 255).astype(np.uint8)

        return images

```

```

[ ]: START_RES = 4
    TARGET_RES = 128

    style_gan = StyleGAN(start_res=START_RES, target_res=TARGET_RES)

```

```

[ ]: def train(
    start_res=START_RES,
    target_res=TARGET_RES,
    steps_per_epoch=5000,
    display_images=True,
):
    opt_cfg = {"learning_rate": 1e-3, "beta_1": 0.0, "beta_2": 0.99, "epsilon": 1e-8}

    val_batch_size = 16
    val_z = tf.random.normal((val_batch_size, style_gan.z_dim))
    val_noise = style_gan.generate_noise(val_batch_size)

    start_res_log2 = int(np.log2(start_res))
    target_res_log2 = int(np.log2(target_res))

```

```

for res_log2 in range(start_res_log2, target_res_log2 + 1):
    res = 2 ** res_log2
    for phase in ["TRANSITION", "STABLE"]:
        if res == start_res and phase == "TRANSITION":
            continue

        train_dl = create_dataloader(res)

        steps = int(train_step_ratio[res_log2] * steps_per_epoch)

        style_gan.compile(
            d_optimizer=tf.keras.optimizers.legacy.Adam(**opt_cfg),
            g_optimizer=tf.keras.optimizers.legacy.Adam(**opt_cfg),
            loss_weights={"gradient_penalty": 10, "drift": 0.001},
            steps_per_epoch=steps,
            res=res,
            phase=phase,
            run_eagerly=False,
        )

        prefix = f"res_{res}x{res}_{style_gan.phase}"

        ckpt_cb = keras.callbacks.ModelCheckpoint(
            f"checkpoints/stylegan_{res}x{res}.ckpt",
            save_weights_only=True,
            verbose=0,
        )
        print(phase)
        style_gan.fit(
            train_dl, epochs=1, steps_per_epoch=steps, callbacks=[ckpt_cb]
        )

        if display_images:
            images = style_gan({"z": val_z, "noise": val_noise, "alpha": 1.
→0})

            plot_images(images, res_log2)

```

```
[ ]: train(start_res=4, target_res=16, steps_per_epoch=1, display_images=True)
```

Model resolution:4x4

STABLE


1/1 [=====] - 9s 9s/step - d_loss: 2.7615 - g_loss:
0.8970



Model resolution:8x8


TRANSITION

1/1 [=====] - 8s 8s/step - d_loss: 5.4979 - g_loss:
1.7689



STABLE


1/1 [=====] - 4s 4s/step - d_loss: 0.6347 - g_loss:
5.9936



Model resolution:16x16

TRANSITION


1/1 [=====] - 10s 10s/step - d_loss: 5.1118 - g_loss:
5.5240



STABLE

WARNING:tensorflow:5 out of the last 5 calls to <function
Model.make_train_function.<locals>.train_function at 0x7ac730eb7910> triggered
tf.function retracing. Tracing is expensive and the excessive number of tracings
could be due to (1) creating @tf.function repeatedly in a loop, (2) passing
tensors with different shapes, (3) passing Python objects instead of tensors.
For (1), please define your @tf.function outside of the loop. For (2),
@tf.function has reduce_retracing=True option that can avoid unnecessary
retracing. For (3), please refer to
https://www.tensorflow.org/guide/function#controlling_retracing and
https://www.tensorflow.org/api_docs/python/tf/function for more details.

1/1 [=====] - 6s 6s/step - d_loss: 0.9506 - g_loss:
11.1828



```
[ ]: url = "https://github.com/soon-yau/stylegan_keras/releases/download/
↳keras_example_v1.0/stylegan_128x128.ckpt.zip"


weights_path = keras.utils.get_file(
    "stylegan_128x128.ckpt.zip",
    url,
    extract=True,
    cache_dir=os.path.abspath("."),
    cache_subdir="pretrained",
)

style_gan.grow_model(128)
style_gan.load_weights(os.path.join("pretrained/stylegan_128x128.ckpt"))

tf.random.set_seed(196)
batch_size = 2
z = tf.random.normal((batch_size, style_gan.z_dim))
w = style_gan.mapping(z)
noise = style_gan.generate_noise(batch_size=batch_size)
images = style_gan({"style_code": w, "noise": noise, "alpha": 1.0})
plot_images(images, 5)
```

Downloading data from https://github.com/soon-yau/stylegan_keras/releases/download/keras_example_v1.0/stylegan_128x128.ckpt.zip
540534982/540534982 [=====] - 11s 0us/step

Model resolution:128x128



```
[ ]: alpha = 0.4
w_mix = np.expand_dims(alpha * w[0] + (1 - alpha) * w[1], 0)
noise_a = [np.expand_dims(n[0], 0) for n in noise]
mix_images = style_gan({"style_code": w_mix, "noise": noise_a})
image_row = np.hstack([images[0], images[1], mix_images[0]])
plt.figure(figsize=(9, 3))
plt.imshow(image_row)
plt.axis("off")
```

```
[ ]: (-0.5, 383.5, 127.5, -0.5)
```

